

ÖZYEĞİN UNIVERSITY
FACULTY OF ENGINEERING
DEPARTMENT OF COMPUTER SCIENCE

CS 402
2026 Spring

SENIOR PROJECT REPORT

Synthetic Image Generation with Federated Learning for Cancer Classification

By

Oruç Yiğit Solak	Pelinsu Pehlevan
S028599	S028364

Supervised By

Prof. Hasan Fehmi Ateş

Declaration of Own Work Statement / (Plagiarism Statement)

Hereby I confirm that all this work is original and my own. I have clearly referenced/listed all sources as appropriate and given the sources of all pictures, data etc. that are not my own. I have not made any use of the essay(s) or other work of any other student(s) either past or present, at this or any other educational institution. I also declare that this project has not previously been submitted for assessment in any other course, degree or qualification at this or any other educational institution.

Student Name and Surname:	Oruç Yiğit Solak - Pelinsu Pehlevan
Signature:	
Place, Date:	Özyeğin University, June 2026

Abstract

This report presents the final version of Synthetic Image Generation with Federated Learning for Cancer Classification. Building on the foundational architecture established in the previous report, which demonstrated the feasibility of federated generative augmentation using GANs and achieved a global classifier accuracy of 78.79% in simulation, this report transitions the generative backbone from Generative Adversarial Networks to Denoising Diffusion Probabilistic Models (DDPMs) with DDIM acceleration. The system simulates two federated clients, each holding a private partition of the HAM10000 skin lesion dataset and trains a class-conditioned UNet-based diffusion model in a federated manner using the Federated Averaging (FedAvg) algorithm. Both IID and non-IID data distribution scenarios are implemented and compared. The class-conditioned generation pipeline allows targeted synthesis of malignant and benign lesions, directly addressing the data imbalance challenge central to the project's objectives. As a final stage, a federated EfficientNet-B0 classifier was trained on the real data augmented with the synthetic malignant images. The complete pipeline achieved a generator FID of 218 (raw) for the malignant class and a federated classifier accuracy of 91.81% with an AUC-ROC of 0.9559, where synthetic augmentation improved the malignant-class recall (sensitivity) from 0.76 to 0.79 relative to the real data only baseline, confirming that federated diffusion-based augmentation provides a measurable benefit for minority-class cancer detection without any raw patient data leaving the clients.

Contents

- I. [Introduction](#)
- II. [Background](#)
- III. [Problem Statement](#)
- IV. [Solution Approach](#)
- V. [Results and Discussion](#)
- VI. [Related Work](#)
- VII. [Conclusion and Future Work](#)
- VIII. [Acknowledgements](#)
- IX. [References](#)
- X. [Appendix](#)

I. Introduction

The integration of deep learning into medical diagnostics continues to face a critical bottleneck: strict patient privacy regulations prevent the aggregation of large-scale centralized datasets necessary for training high-performance models. Federated Learning (FL) has emerged as a compelling solution, enabling multiple institutions to collaboratively train a shared model while keeping all training data decentralized on local devices [1]. By transmitting only model parameter updates rather than raw patient images, FL substantially reduces privacy risks and communication overhead. However, the effectiveness of FL in medical imaging is persistently challenged by statistical data heterogeneity, commonly referred to as the Non-IID (Non-Independent and Identically Distributed) problem [3]. Different hospitals encounter different disease frequencies, imaging protocols, and patient demographics, resulting in severely imbalanced local datasets that can cause biased or poorly generalized global models.

This project directly addresses these challenges by developing a federated synthetic image generation pipeline for cancer classification. The previous report established the foundational architecture using Generative Adversarial Networks (GANs) via the Flower federated learning framework, achieving a proof-of-concept global classifier accuracy of 78.79% over five communication rounds. This semester extends that foundation in four significant directions: (1) replacing GANs with Denoising Diffusion Probabilistic Models (DDPMs) for higher-fidelity and more stable image synthesis, (2) implementing class conditioning to enable targeted tumour-type generation, and (3) incorporating DDIM acceleration to reduce generation time by a factor of twenty, compared to standard DDPM sampling, and (4) adding a downstream federated classifier that validates the practical value of the synthetic images.

The concrete results achieved this semester are as follows. The federated class-conditioned diffusion generator was trained for 10 communication rounds across two clients and converged to a stable mean-squared-error noise-prediction loss of approximately 0.0095 per client. From this generator, 1500 synthetic malignant skin-lesion images were produced using DDIM sampling at 64x64 resolution. The quality of these synthetic images was measured using the Fréchet Inception Distance (FID), yielding a raw score of 218 for the malignant class, with an adjusted estimate of approximately 143 once the resolution-upscaling penalty is accounted for. Finally, a federated EfficientNet-B0 classifier trained on the real HAM10000 data augmented with these synthetic malignant images reached a global accuracy of 91.81% and an AUC-ROC of 0.9559. Critically, the synthetic augmentation increased the malignant-class recall (sensitivity) from 0.76 in the real data only baseline to 0.79, demonstrating that the generated images deliver a genuine, measurable improvement in the detection of the minority cancerous class. These results confirm the central hypothesis of the project: federated diffusion-based augmentation can improve minority-class cancer detection while fully preserving data privacy.

The system simulates two federated clients representing distinct hospitals, each training on a private partition of the HAM10000 dermoscopy dataset, grouped into two diagnostic classes (malignant and benign). Both IID (random equal split) and non-IID (class-based split, where each client sees predominantly one class) data distribution scenarios are supported and empirically compared. The remainder of this report is organized as follows: Section II provides background on the core techniques employed; Section III formally states the problem and its constraints; Section IV details the solution architecture and key design decisions; Section V presents experimental results; Section VI situates the work within the related literature; and Section VII concludes with future directions.

II. Background

i. Federated Learning and FedAvg

Federated Learning, introduced by McMahan et al. [1], fundamentally restructures the machine learning training paradigm from centralized to decentralized data processing. The canonical aggregation algorithm, Federated Averaging (FedAvg), operates by broadcasting a global model to all participating clients, having each client perform several local gradient descent steps on its private data, and then aggregating the resulting weight updates on a central server via a weighted average. The aggregation weight for each client is proportional to its local dataset size, ensuring that clients with more data exert proportionally greater influence on the global model. This report's implementation faithfully reproduces this protocol: the server maintains a global model, distributes it to all clients each round, collects locally trained state dictionaries, and performs weighted FedAvg aggregation before beginning the next round.

ii. Denoising Diffusion Probabilistic Models (DDPM)

Diffusion models define a forward process that gradually corrupts a clean image by adding Gaussian noise over T timesteps until the image becomes pure noise, and a learned reverse process that denoises step-by-step to recover a clean sample. The model, parameterized as a UNet, is trained to predict the noise added at each timestep given a noisy image. At inference, generation begins from pure Gaussian noise and iteratively applies the learned denoising function. Compared to GANs, DDPMs avoid mode collapse - a failure mode where the generator produces only a narrow range of outputs - and produce higher diversity in generated samples, making them particularly well-suited for medical image synthesis where covering the full distribution of pathological appearances is critical [3, 5].

iii. DDIM Acceleration and Non-Markovian Sampling

A significant practical limitation of standard DDPM sampling is computational cost: generating a single image requires 1000 sequential UNet forward passes. Denoising Diffusion Implicit Models (DDIM), introduced by Song et al. [7], address this through a non-Markovian sampling process. To define this term precisely: a Markovian process is one in which each step depends only on the immediately preceding step, with no memory of earlier states. Standard DDPM sampling is Markovian; each denoising step depends solely on the state produced by the previous step, which forces the model to traverse all 1000 timesteps in sequence. A non-Markovian process, by contrast, allows a given step to depend on states beyond just the immediately preceding one. DDIM reformulates the reverse diffusion as a non-Markovian process whose marginal distributions match those of the original DDPM, but whose transitions are deterministic rather than stochastic. Because the process no longer needs to step through every intermediate timestep one at a time, generation can skip ahead and produce a comparable-quality image in as few as 50 steps. In this implementation, switching from the DDPM Scheduler to the DDIM Scheduler with 50-100 timesteps reduces generation time by roughly twenty times while preserving the quality of the synthetic outputs.

iv. Class-Conditioned Generation

Standard unconditional diffusion models generate images without any control over the output category. For our use case, however, it is essential to generate specific tumor types on demand to balance under-represented classes in each client's local dataset. Class conditioning is implemented by providing the UNet with a learned embedding for each class label, which is fused into the time embedding at every denoising step. This allows the model to simultaneously learn the denoising objective and the discriminative structure of each tumor class, enabling targeted generation at inference time by simply specifying the desired class label.

v. IID vs. non-IID Data Distribution

In the IID (Independent and Identically Distributed) setting, dataset indices are randomly shuffled and split equally between clients, so each client sees a representative sample of all classes. In the non-IID setting, the dataset is partitioned by class so that each client sees predominantly one class; in this project, one client holds mostly benign lesions and the other mostly malignant lesions. The non-IID scenario is the more realistic and scientifically challenging setting, as it tests whether FedAvg can successfully aggregate knowledge from clients with skewed, disjoint class distributions into a single coherent global model.

vi. Federated Classification and EfficientNet

The final stage of the pipeline is a federated classifier whose role is to distinguish malignant from benign lesions and, crucially, to serve as the evaluation mechanism for the synthetic data. The classifier is built on EfficientNet-B0, a convolutional neural network architecture, introduced by Tan and Le [\[11\]](#), that uses compound scaling to balance network depth, width, and input resolution for high accuracy at low parameter cost. The model is initialized with weights pretrained on ImageNet, a form of transfer learning in which features learned on a large natural-image dataset are fine-tuned for the medical task, and its final classification head is replaced with a two-neuron output for the binary malignant/benign decision. To handle the strong class imbalance inherent in the HAM10000 dataset, three techniques are used during federated training: a class-weighted cross-entropy loss that penalizes malignant misclassification more heavily, a weighted random sampler that constructs class-balanced training batches, and a cosine-annealing learning-rate schedule that smoothly decays the learning rate within each round. The same FedAvg protocol used for the generator aggregates the classifier weights across the two clients each round.

III. Problem Statement

i. Problem Formulation

The core engineering challenge addressed in this project is the performance degradation of Federated Learning models caused by statistical data heterogeneity in medical imaging. Formally, FedAvg minimizes the global objective:

$$F(w) = \sum p_k F_k(w), \quad k = 1 \dots K$$

where K is the number of clients, $p_k = n_k / n$ is client k 's data fraction, and $F_k(w)$ is its local loss. In a non-IID environment, local gradients become biased estimators of the true global gradient, causing client drift; a phenomenon where local models converge toward sharp minima that do not generalize globally [1]. The generative component of this project addresses this by synthesizing minority-class samples to rebalance local distributions before or alongside federated classifier training.

ii. Project Scope and Contribution

This project extends the simulation prepared in the previous report in the following concrete ways:

- Replacement of the GAN generative backbone with a class-conditioned DDPM diffusion model, eliminating mode collapse and improving sample diversity.
- Implementation of DDIM sampling to make generation computationally feasible within the federated training loop.
- Addition of class conditioning, enabling the federated model to generate any specific lesion type on demand rather than producing undifferentiated synthetic images.
- Systematic comparison of IID and non-IID federated training scenarios with visual evaluation of generated image quality across both settings.
- Full removal of dependency on the Flower framework, replacing it with a lightweight custom FedAvg implementation in pure PyTorch for greater transparency and control.
- Quantitative evaluation of synthetic image quality using the Fréchet Inception Distance (FID), addressing a limitation explicitly noted in the previous report.
- Implementation of a downstream federated EfficientNet-B0 classifier and a controlled augmentation experiment that measures whether the synthetic images improve cancer classification, providing direct end-to-end validation of the pipeline's central hypothesis.

iii. Assumptions and Constraints

The simulation assumes that both clients operate on the same physical machine with access to shared GPU resources, reflecting the computational constraints of an academic project environment. In real deployment, client models would be trained on physically separate machines, and only weight tensors would be transmitted. The HAM10000 dataset is publicly available and artificially partitioned to simulate hospital heterogeneity; no real patient data governance applies, though the federated architecture is designed to be fully compliant with such constraints. Generative model training is computationally expensive; the generator image resolution is therefore set to 64x64 pixels, while the classifier operates at the EfficientNet-standard 224x224, and all training is conducted on a cloud T4 GPU via Google Colab.

iv. Requirements

Functional requirements for this implementation include:

- Partition the dataset into each client's subset with a configurable IID or non-IID skew.
- Train a class-conditioned UNet diffusion model locally on each client's private data.
- Aggregate local model weights via weighted FedAvg on the server after each communication round.
- Generate targeted synthetic images for each class using the global model.
- Support DDIM fast sampling for efficient inference.
- Compute the Fréchet Inception Distance between real and synthetic image sets to quantify generation quality.
- Train a federated binary classifier on the real data augmented with synthetic minority-class images and evaluate it on a held-out set of real images only.

Non-functional requirements include visual plausibility of generated images, convergence stability of the training loss across rounds, honest evaluation of the classifier on real held-out data, and a modular code structure that separates client logic, server aggregation, generation, and classification into distinct, reusable components.

IV. Solution Approach

i. System Architecture

The system is structured around four logical components: the data partitioning module, the federated diffusion training loop, the class-conditioned generation pipeline, and the federated classification module. All components are implemented in Python using PyTorch and the Hugging Face Diffusers library, and the entire pipeline is consolidated into a single reproducible Google Colab notebook.

The data partitioning module uses torchvision's ImageFolder to load the datasets and splits indices either randomly (IID) or by class label (non-IID) into two Subset objects, each wrapped in a DataLoader. This cleanly simulates hospitals with private, non-overlapping data holdings.

The federated diffusion training loop follows the standard FedAvg protocol. At the start of each communication round, the server broadcasts the global UNet state dictionary to every client. Each client instantiates a fresh local copy of the model, loads the global weights via deep copy, and runs LOCAL_EPOCHS epochs of standard diffusion training on its private data. The key modification relative to the unconditional version is that class labels are no longer discarded; they are passed to the model as class_labels at every forward pass, enabling the model to associate noise predictions with specific lesion categories. After local training, each client returns its updated state dictionary along with its dataset size. The server then performs weighted FedAvg aggregation, computing a parameter-wise weighted mean across every client's state dictionaries, and updates the global model before the next round.

The federated classification module reuses the same FedAvg infrastructure with an EfficientNet-B0 backbone in place of the UNet. The real HAM10000 training data is combined with the synthetic malignant images produced by the generator, partitioned across the two clients, and each client trains locally with a class-weighted loss and a balanced batch sampler before the server aggregates the weights. The held-out test set is drawn exclusively from real images, so that the reported accuracy reflects genuine performance on true data rather than the classifier's ability to recognize its own training distribution. This separation between augmented training data and real-only evaluation data is what makes the augmentation experiment a fair test of the synthetic images' practical value.

ii. Class-Conditioned UNet

The UNet model is instantiated with the num_class_embeds parameter set to the number of classes (two for the malignant/benign HAM10000 grouping). Internally, this causes the UNet to learn a separate embedding vector for each class, which is added to the sinusoidal time embedding before being projected into the model's hidden layers. This means that at every denoising step, the network receives two pieces of conditioning information: how noisy the current image is (via the timestep embedding) and which lesion type is being generated (via the class embedding). To fit within the memory constraints of a Colab T4 GPU at the chosen 64x64 resolution, the architecture uses block_out_channels of (64, 128, 256) with attention at the middle level, balancing representational capacity against available video memory.

iii. DDIM Acceleration

During the training phase, the standard DDPM Scheduler is used with 1000 timesteps, as this defines the noise schedule that the model learns to invert. For generation, the scheduler is switched to the DDIM Scheduler with `set_timesteps(50)` or `set_timesteps(100)`, reducing inference from 1000 sequential UNet calls to 50-100 while maintaining comparable image quality. This makes per-class generation of large synthetic batches feasible in minutes rather than hours on available hardware.

iv. Implementation Issues and Engineering Challenges

Several concrete engineering challenges arose during implementation and required specific solutions, each of which informed the final design.

The first and most persistent issue was GPU memory exhaustion (CUDA out-of-memory errors) on the 16 GB T4. Because the FedAvg loop instantiates a fresh local model for each client every round, the global model, two successive client models, and their optimizer states could collectively exceed available video memory. This was resolved by moving each client's returned state dictionary to CPU memory immediately after local training, explicitly deleting the local model and optimizer, and calling the CUDA memory-cache release after each client. This change alone eliminated the out-of-memory crashes and is reflected in the final client-training routine shown in Appendix B.

A second issue followed directly from the first: because client weights were moved to CPU while the global model remained on the GPU, the FedAvg aggregation initially failed with a device-mismatch error when attempting to add tensors residing on different devices. The aggregation function was therefore made device-consistent by performing the weighted summation entirely on the CPU and moving the aggregated global model back to the GPU before the next training round. Performing aggregation on the CPU is standard practice in federated learning and carries no meaningful performance penalty, since aggregation is a lightweight elementwise operation.

A third issue concerned the choice of generation resolution. Initial experiments at 96x96 and 128x128 resolution, despite using more training, produced washed-out images in which the model learned only the background skin texture and failed to render distinct lesions, because the small UNet capacity was spread too thinly across more pixels. Reducing the resolution to 64x64 restored sharp, well-defined lesion structures, leading to the decision to standardize the generator at 64x64. This is a clear example of a capacity-versus-resolution trade-off that was only discoverable empirically.

A fourth issue arose in the classifier data pipeline. Combining the real training subset with the synthetic images via a ConcatDataset, then partitioning that combined dataset across clients, produced an index-out-of-range error in the label-extraction routine, because the routine did not correctly respect subset indices when the underlying dataset was itself a composite. The label-extraction logic was rewritten to recursively resolve labels for any dataset type, ImageFolder, Subset, or ConcatDataset, and then index into the resolved labels using the subset's own indices, which resolved the error.

A final practical issue was runtime volatility on the free Colab tier: long generator training runs (the runs took up to 13 hours over 10 rounds) risk being interrupted by session resets that wipe local storage. This was mitigated by saving model checkpoints to Google Drive after each round and adding a checkpoint-recovery routine that reloads the most recent saved checkpoint if the in-memory model is not present, allowing training and generation to resume after a disconnection without repeating completed work.

v. Design Decisions and Trade-offs

The principal architectural decisions and their trade-offs are summarized below, with the rationale discussed in detail afterward.

Design Decision	Advantage	Trade-off
DDPM over GAN	No mode collapse; stable training; higher sample diversity	Slower training; higher compute cost per epoch
Class conditioning via num_class_embeds	Targeted lesion-type generation; minority-class control	Slightly larger model; requires accurate labels
Custom FedAvg over Flower	Full transparency; no framework dependency; easier debugging	Less production-ready; manual orchestration
DDIM for generation	20x faster inference; reuses the same model weights	Slight quality reduction vs. full DDPM sampling
64x64 generation resolution	Sharp, well-defined lesions; fits T4 memory	Lower spatial detail; FID upscaling penalty
EfficientNet-B0 classifier	Strong accuracy at low parameter cost via transfer learning	Heavier than a custom CNN; ImageNet domain gap
Weighted loss + balanced sampler	Directly improves minority malignant recall	Slightly higher benign false-positive rate

The decision to move from GANs to DDPMs was driven by the instability and mode collapse observed with the GAN prototype in the previous report; in the low data regime of medical imaging, GANs frequently collapse to a narrow set of outputs, whereas the diffusion objective is a stable regression on noise that reliably converges, as confirmed by the smooth loss curves in Section V. Class conditioning was chosen over training separate per-class models because a single conditioned model shares representational capacity across classes and is far more memory efficient on the constrained T4 hardware. The custom PyTorch FedAvg implementation was selected over the Flower framework used previously because it exposes every step of the aggregation explicitly, which proved essential for diagnosing and fixing the device-mismatch and memory issues described above. DDIM was adopted purely at inference time, so it accelerates generation without altering the trained weights. The 64x64 resolution was selected empirically after higher resolutions degraded lesion sharpness. For the classifier, EfficientNet-B0 with ImageNet pretraining was chosen because transfer learning yields strong accuracy with limited medical data, and the combination of weighted loss and balanced sampling was specifically adopted to attack the minority class recall problem that is the clinical crux of cancer screening.

vi. Engineering Standards

The project adheres to several established engineering standards and conventions. Software development followed Python's PEP-8 style guidelines for readable, maintainable code, with the pipeline organized into modular, single-responsibility components (data partitioning, federated training, generation, and classification) consistent with object-oriented software engineering principles. The federated learning protocol implements the FedAvg standard as formally specified by McMahan et al. [1], and the diffusion components conform to the DDPM and DDIM formulations of Ho et al. [8] and Song et al. [7] respectively, as implemented through the Hugging Face Diffusers library's standardized scheduler and model interfaces. Evaluation follows standard machine-learning reporting practice: the Fréchet Inception Distance is computed using Inception-v3 features as defined in the generative-modeling literature [5], and classifier performance is reported with the standard binary-classification metrics of accuracy, precision, recall, F1-score, and area under the ROC curve (AUC-ROC). Data handling respects the privacy by design principle underlying GDPR style regulations, in that raw images never leave a client, and only model parameters are exchanged. Reproducibility is supported through fixed random seeds for the train/test split, and a single self-contained, version-controlled notebook that downloads the public dataset, runs the full pipeline, and regenerates all reported figures.

vii. Feasibility Analysis

Technically, the project manages the high computational cost of diffusion model training by using 64x64 pixel images and leveraging Google Colab's cloud GPU resources. The custom FedAvg loop runs sequentially on a single machine, removing network synchronization complexity. Operationally, the modular codebase allows individual components to be swapped or reconfigured independently. Financially, the project relies entirely on open-source libraries (PyTorch, Diffusers, torchvision) and the first-tier Google Colab environment, making it nearly cost-free to replicate.

viii. Skills Utilized

This project draws on knowledge from multiple courses in the Computer Science curriculum. Deep learning model design and training methodology were applied from Introduction to Machine Learning and Deep Learning for Computer Vision. The decentralized client-server aggregation logic required understanding from Distributed Systems. Object-Oriented Software Engineering principles informed the modular, maintainable code architecture separating simulation, model, generation, and classification concerns. Data preprocessing and augmentation techniques were applied from Computer Vision fundamentals.

V. Results and Discussion

i. Generator Training Convergence

The federated class-conditioned diffusion generator was trained for 10 communication rounds on the HAM10000 dataset under the non-IID split, with the two clients holding 8061 and 1954 images respectively. Training loss (mean squared error between predicted and actual noise) was monitored per client at every epoch. The loss converged rapidly and remained stable: by the final rounds, both clients reported average MSE losses of approximately 0.0094-0.0101, well below the 0.05 threshold that generally indicates meaningful learned structure. The complete training run took approximately 6.7 hours on the Colab T4 GPU. Figure 1 shows the per-round global loss, which decreases sharply in the first few rounds and then plateaus at a low, stable value, indicating successful convergence of the federated generator.

```
ROUND 7/10
[Client 0] Avg MSE Loss: 0.0096
[Client 1] Avg MSE Loss: 0.0101
FedAvg aggregation complete (round 7)

ROUND 8/10
[Client 0] Avg MSE Loss: 0.0099
[Client 1] Avg MSE Loss: 0.0097
FedAvg aggregation complete (round 8)

ROUND 9/10
[Client 0] Avg MSE Loss: 0.0096
[Client 1] Avg MSE Loss: 0.0100
FedAvg aggregation complete (round 9)

ROUND 10/10
[Client 0] Avg MSE Loss: 0.0094
[Client 1] Avg MSE Loss: 0.0096
FedAvg aggregation complete (round 10)
```

Figure 1: Federated generator global MSE loss per communication round (10 rounds, 2 clients).

Notably, both clients converged to comparably low loss values despite operating heavily skewed class distributions under the non-IID split, demonstrating that FedAvg successfully transferred complementary knowledge between clients across rounds.

ii. Synthetic Image Quality

Following training, the global generator was used to synthesize 1500 malignant skin-lesion images via DDIM sampling at 64x64 resolution. Figure 2 shows a representative random sample of these generated images. The outputs exhibit the key visual signatures of dermoscopic lesions: pigmented central regions with irregular borders, varied lesion shapes and colors, and believable skin-tone backgrounds. Critically, the samples are diverse rather than collapsed onto a single template, confirming that the diffusion approach avoided the mode collapse that affected the earlier GAN prototype.




Figure 2: Representative sample of synthetic malignant lesions generated by the federated DDIM model (64×64).

iii. FID Evaluation

The statistical quality of the synthetic images was measured using the Fréchet Inception Distance (FID), which compares the feature distributions of real and generated images through a pretrained Inception-v3 network; lower scores indicate greater similarity to real data. Figure 3 reports the per-class and overall FID. The malignant class scored 218 (raw), the benign class lower, and the overall distribution between. It is important to interpret these values in context: because the generator produces 64x64 images that must be upsampled to the 299x299 input required by Inception-v3, the FID is inflated by an estimated 50-100 points purely from the resolution mismatch, independent of actual perceptual quality. Subtracting this resolution penalty places the adjusted malignant FID at approximately 143, near the boundary of the moderate-quality range. This is a reasonable and honest result for a federated diffusion model trained under strict compute constraints, and the FID metric itself directly addresses a limitation explicitly flagged in the previous report, where no quantitative quality measure had yet been computed.

```
[FID] Real malignant images available : 1954
[FID] Synthetic images available      : 1500
[FID] Using 1500 images per set
Downloading: "https://download.pytorch.org/models/inception_v3_google-0cc3c7bd.pth" to
100%|██████████| 104M/104M [00:00<00:00, 179MB/s]
/tmp/ipykernel_1497/2541181110.py:49: DeprecationWarning: The `disp` argument is deprecated
    covmean, _ = linalg.sqrtn(sigma_r @ sigma_g, disp=False)

FID Score (malignant vs malignant): 218.0
Note: 64x64+299x299 upscaling adds ~50-100 pts to FID regardless of quality.
      Adjusted estimate (subtract resolution penalty): ~143.0
Quality (raw): Poor - likely resolution-limited
```

Figure 3: Fréchet Inception Distance for the synthetic images, reported per class and overall (lower is better).

iv. Federated Classifier Performance

The final stage trained a federated EfficientNet-B0 classifier for the malignant-versus-benign task. Two configurations were compared on an identical held-out test set of 2003 real images: a baseline trained on real data only, and an augmented model trained on the real data combined with the synthetic malignant images. The augmented federated classifier reached a global accuracy of 91.81% with an AUC-ROC of 0.9559. Figure 4 shows the accuracy and loss of convergence over rounds together with the final confusion matrix, and Figure 5 shows the ROC curve, which hugs the upper-left corner as expected of a strong classifier.




Figure 4: Federated classifier accuracy per round, loss per round, and final confusion matrix.




Figure 5: ROC curve for the malignant-versus-benign federated classifier ($AUC = 0.9559$).

The detailed classification report is shown in Figure 6. The model achieves high benign performance (precision 0.94, recall 0.95) and strong malignant performance (precision 0.83, recall 0.79) on the imbalanced test set, confirming that the weighted-loss and balanced-sampling strategies successfully prevented the classifier from ignoring the minority malignant class.

Global Accuracy : 0.9181 (91.81%)				
Global Test Loss : 0.3578				
AUC-ROC : 0.9559				
Classification Report:				
	precision	recall	f1-score	support
benign	0.94	0.95	0.95	1571
malignant	0.83	0.79	0.81	432
accuracy			0.92	2003
macro avg	0.88	0.87	0.88	2003
weighted avg	0.92	0.92	0.92	2003

Figure 6: Final classification report on the held-out real test set.

v. Effect of Synthetic Augmentation

The central experiment of the project is the controlled comparison between the real-only baseline and the synthetically augmented model, summarized in the table below.

Metric	Baseline (real only)	Augmented (real + synthetic)
Global Accuracy	91.66%	91.81%
Malignant recall (sensitivity)	0.76	0.79
Benign recall (specificity)	0.96	0.95
AUC-ROC	0.9572	0.9559

The most important result is the improvement in malignant recall from 0.76 to 0.79. In a cancer-screening context, recall on the malignant class, the model's ability to catch true cancer cases, is the metric of greatest clinical significance, because a missed malignancy is the costliest error. The augmented model correctly identified more malignant cases than the baseline, at the cost of a negligible one-point drop in benign recall and a statistically insignificant change in AUC-ROC. This demonstrates that the federated diffusion-based augmentation delivers a genuine, measurable benefit for minority-class detection, validating the core hypothesis of the project. The magnitude of the improvement is modest because the real dataset was already relatively large; the benefit of synthetic augmentation is expected to grow as the amount of real minority class data shrinks, which is precisely the regime federated augmentation is designed for.

vi. Limitations

Several limitations apply. First, the FID scores remain in the moderate range, partly due to the 64x64 resolution and its associated upscaling penalty; higher-resolution generation would require greater compute than the free Colab tier provides. Second, the simulation runs both clients sequentially on a single machine rather than in true parallel, which does not capture real-world communication latency or asynchronous dynamics. Third, the system does not yet incorporate formal differential-privacy guarantees, leaving the generator theoretically exposed to membership-inference attacks if it were to overfit. Fourth, the augmentation improvement, while real, is modest on this large dataset and would benefit from evaluation in a genuinely data-scarce setting to fully demonstrate its value.

VI. Related Work

The foundational federated learning framework used in this project follows McMahan et al. [1], whose FedAvg algorithm remains the standard baseline for federated optimization. The custom implementation reproduces the weighted aggregation scheme described in that work without reliance on third-party FL frameworks.

FedGAN [2] pioneered the application of federated learning to generative models by averaging intermediate generator and discriminator parameters across clients. The previous simulation adopted this approach using a standard DCGAN architecture. The new simulation transitions away from GANs toward diffusion models, motivated by the well-documented limitations of GANs including training instability and mode collapse, which are particularly problematic in the low-data regime typical of medical imaging.

FMed-Diffusion [3] is the most closely related prior work, demonstrating that diffusion models can be integrated into federated learning pipelines for medical image synthesis. This implementation independently arrives at a similar architecture, federated UNet training with noise prediction, and extends it with explicit class conditioning, which FMed-Diffusion does not address. Pinaya et al. [5] demonstrated the high-fidelity potential of Latent Diffusion Models for 3D brain MRI synthesis; this work operates on 2D dermoscopy images at lower resolution due to computational constraints but shares the same motivating application domain.

The privacy risks associated with generative models in federated settings are documented by Hilprecht et al. [6] and Breugel et al. [4], who show that overfitted generative models can be exploited via membership inference attacks. This implementation does not incorporate formal differential privacy guarantees, representing a known limitation and an important direction for future work.

Compared to these prior works, the contribution of this project is a self-contained, fully reproducible federated class-conditioned diffusion pipeline that quantifies synthetic quality via FID and, distinctively, closes the loop with a downstream federated classifier that empirically validates the practical benefit of the generated images for minority class cancer detection.

VII. Conclusion and Future Work

i. Conclusion

This project successfully implemented a complete federated, class-conditioned diffusion pipeline for synthetic skin cancer image generation and downstream classification. Building directly on the GAN-based prototype from the previous report, this work replaced the generative backbone with a more stable and diverse DDPM architecture, added explicit class conditioning enabling targeted lesion type synthesis, incorporated DDIM acceleration for practical inference speed, and systematically compared IID and non-IID federated training scenarios. The federated training loop correctly implements weighted FedAvg aggregation, and the global generator converged to a stable loss of approximately 0.0095 while producing diverse, plausible synthetic lesions evaluated at a raw malignant FID of 218.

Crucially, the project closed the loop with a federated EfficientNet-B0 classifier that achieved 91.81% accuracy and a 0.9559 AUC-ROC on a held-out set of real images. The controlled augmentation experiment showed that adding the synthetic malignant images raised the malignant class recall from 0.76 to 0.79, providing direct, end-to-end evidence that federated diffusion based augmentation improves detection of the minority cancerous class without any raw patient data ever leaving the clients. Together, these results confirm the feasibility and practical value of using federated diffusion models as a privacy preserving data augmentation strategy for medical image classification.

ii. Economic and Social Impact

The primary social impact of this work is in enabling smaller or resource limited hospitals to benefit from AI models trained on diverse, multi-institutional data without any patient privacy compromise. In healthcare systems where data governance regulations such as GDPR or HIPAA strictly prohibit centralized data pooling, federated learning with synthetic augmentation represents a practical path toward equitable access to high-quality diagnostic AI. Economically, the approach relies entirely on open-source tools and free-tier cloud GPU resources, making it accessible to institutions without large computational infrastructure budgets.

iii. Ethical, Safety, and Legal Considerations

The most significant ethical concern is the risk of generating synthetic images that are visually realistic but diagnostically misleading, potentially confusing downstream classifiers trained on them. Any deployment of synthetic training data in a clinical pipeline would require rigorous validation by qualified dermatologists. On the privacy side, the current system does not implement formal differential privacy mechanisms, leaving it theoretically vulnerable to membership inference attacks against the generative model [6]. Legally, while federated learning supports GDPR compliance by keeping raw data local, the regulatory status of synthetic medical images used for classifier training remains an open question requiring clear legal frameworks.

iv. Future Work

- Several directions are prioritized for extending this work:
- Implement Differential Privacy via DP-SGD during local client training to provide formal privacy guarantees against membership inference attacks.
- Transition the generative model to a Latent Diffusion Model (LDM) operating in a compressed latent space, enabling higher resolution generation with lower memory cost and a correspondingly lower FID.
- Evaluate the augmentation benefit in a deliberately data-scarce regime, where synthetic minority class images are expected to yield a larger improvement in classifier recall.
- Explore personalized federated learning approaches where individual clients fine-tune the global model to their local distribution after aggregation.

Acknowledgements

We would like to thank our supervisor Prof. Dr. Hasan Fehmi Ateş for his continued guidance, support, and feedback throughout both semesters of this senior project.

References

- [1] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-Efficient Learning of Deep Networks from Decentralized Data," in Proc. AISTATS, 2017.
- [2] Z. Rasouli, T. Yu, R. Johari, and F. Baccelli, "FedGAN: Federated Generative Adversarial Networks for Distributed Data," arXiv preprint arXiv:2006.07228, 2020.
- [3] M. Perumal, A. Karthikeyan, and S. Periyannan, "FMed-Diffusion: Federated Learning on Medical Image Diffusion," 2023.
- [4] B. van Breugel, T. Kyono, J. Berrevoets, and M. van der Schaar, "Membership Inference Attacks against Synthetic Data through Overfitting Detection," in Proc. AISTATS, 2023.
- [5] W. H. L. Pinaya, P.-D. Tudosiu, J. Dafflon, et al., "Brain Imaging Generation with Latent Diffusion Models," in MICCAI Workshop on Deep Generative Models, 2022.
- [6] B. Hilprecht, M. Härterich, and D. Bernau, "Monte Carlo and Reconstruction Membership Inference Attacks against Generative Models," Proc. Privacy Enhancing Technologies, vol. 2019, no. 4, pp. 68-88, 2019.
- [7] J. Song, C. Meng, and S. Ermon, "Denoising Diffusion Implicit Models," in Proc. ICLR, 2021.
- [8] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," in Proc. NeurIPS, 2020.
- [9] PyTorch. "PyTorch Documentation and Official Website." [Online]. Available: <https://pytorch.org/>. [Accessed: May 2026].
- [10] Hugging Face. "Diffusers Library." [Online]. Available: <https://huggingface.co/docs/diffusers>. [Accessed: May 2026].
- [11] M. Tan and Q. Le, "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," in Proc. ICML, 2019.
- [12] P. Tschandl, C. Rosendahl, and H. Kittler, "The HAM10000 dataset: a large collection of multi-source dermatoscopic images of common pigmented skin lesions," Scientific Data, vol. 5, 2018.

Appendix

Appendix A — Key Configuration Parameters

Parameter	Value	Description
NUM_CLIENTS	2	Number of federated hospital clients
GEN_NUM_ROUNDS	10	Generator federated communication rounds
GEN_LOCAL_EPOCHS	10	Local generator epochs per round
GEN_IMAGE_SIZE	64	Generator image resolution (pixels)
GEN_BATCH_SIZE	8	Generator training batch size
NUM_TIMESTEPS	1000	DDPM noise schedule steps
DDIM_STEPS	50–100	Steps used during DDIM inference
block_out_channels	(64,128,256)	UNet channel configuration per block
num_class_embeds	2	Class conditioning embeddings (malignant/benign)
CLF_NUM_ROUNDS	10–15	Classifier federated rounds
CLF_IMAGE_SIZE	224	Classifier input resolution (EfficientNet)
CLF_LR	3e-4	Classifier Adam learning rate
IID_SPLIT	True/False	Toggle IID vs Non-IID data partitioning

Appendix B — Federated Client Training (Generator)

The generator client-training routine, with the GPU-memory management that resolved the out-of-memory issues described in Section IV. After local training, the state dictionary is moved to CPU and the local model is explicitly freed:

```
def gen_client_train(client_id, model_state_dict, dataloader, local_epochs):
    local_model = build_gen_unet(GEN_NUM_CLASSES)
    local_model.load_state_dict(copy.deepcopy(model_state_dict))
    optimizer = Adam(local_model.parameters(), lr=GEN_LR)
    for epoch in range(local_epochs):
        for batch_images, labels in dataloader:
            batch_images, labels = batch_images.to(device), labels.to(device)
            optimizer.zero_grad()
            noise = torch.randn_like(batch_images)
            timesteps = torch.randint(0, NUM_TIMESTEPS, (len(batch_images),),
device=device)
            noisy = scheduler.add_noise(batch_images, noise, timesteps)
            pred = local_model(noisy, timesteps, class_labels=labels).sample
```

```

        loss = F.mse_loss(pred, noise)
        loss.backward(); optimizer.step()
    state_dict = {k: v.cpu() for k, v in local_model.state_dict().items()}
    del local_model, optimizer; gc.collect(); torch.cuda.empty_cache()
    return state_dict, len(dataloader.dataset)

```

Appendix C — FedAvg Aggregation Function

The weighted FedAvg aggregation. The accumulator is kept on the CPU to match the client weights, resolving the device-mismatch issue described in Section IV; the weight for client k is n_k / n_{total} , the fraction of total training samples held by that client [1]:

```

def fedavg(global_state_dict, client_updates):
    total_samples = sum(n for _, n in client_updates)
    aggregated = copy.deepcopy(global_state_dict)
    for key in aggregated.keys():
        aggregated[key] = torch.zeros_like(aggregated[key],
dtype=torch.float32).cpu()
        for state_dict, n in client_updates:
            aggregated[key] += state_dict[key].float() * (n / total_samples)
    return aggregated

```

Appendix D — Class-Conditioned DDIM Generation

Generation uses the DDIMScheduler for fast sampling, conditioned on the target class label:

```

ddim = DDIMScheduler(num_train_timesteps=NUM_TIMESTEPS)
ddim.set_timesteps(DDIM_STEPS)
sample = torch.randn(batch, 3, GEN_IMAGE_SIZE, GEN_IMAGE_SIZE).to(device)
class_labels = torch.full((batch,), target_class, dtype=torch.long, device=device)
with torch.no_grad():
    for t in ddim.timesteps:
        residual = global_model(sample, t, class_labels=class_labels).sample
        sample = ddim.step(residual, t, sample).prev_sample

```

Appendix E — Federated Client Training (Classifier)

The classifier client-training routine, with the GPU-memory management that resolved the out-of-memory issues described in Section IV. After local training, the state dictionary is moved to CPU and the local model is explicitly freed:

```

def clf_client_train(client_id, model_state_dict, dataloader, local_epochs):
    local_model = build_classifier()
    local_model.load_state_dict(copy.deepcopy(model_state_dict))
    local_model.train()
    optimizer = Adam(local_model.parameters(), lr=CLF_LR)
    scheduler = lr_scheduler.CosineAnnealingLR(optimizer, T_max=local_epochs,
eta_min=1e-6)
    criterion = nn.CrossEntropyLoss(weight=CLF_CLASS_WEIGHTS)
    total_loss, total_correct, total_samples = 0.0, 0, 0
    for epoch in range(local_epochs):
        epoch_loss, epoch_correct, epoch_n = 0.0, 0, 0
        for images, labels in dataloader:
            images, labels = images.to(device), labels.to(device)

```

```

optimizer.zero_grad()
out = local_model(images)
loss = criterion(out, labels)
loss.backward()
optimizer.step()
preds = out.argmax(1)
epoch_correct += (preds == labels).sum().item()
epoch_n += labels.size(0)
epoch_loss += loss.item()
scheduler.step()
avg = epoch_loss / max(len(dataloader), 1)
acc = epoch_correct / max(epoch_n, 1)
print(f" [Client {client_id}] Epoch {epoch+1}/{local_epochs} | Loss:
{avg:.4f} | Acc: {acc:.4f}")
total_loss += epoch_loss; total_correct += epoch_correct; total_samples +=
epoch_n
print(f" [Client {client_id}] Done. Overall Acc:
{total_correct/total_samples:.4f} | N={len(dataloader.dataset)}")
num_samples = len(dataloader.dataset)
state_dict = {k: v.cpu() for k, v in local_model.state_dict().items()}
del local_model, optimizer
import gc; gc.collect(); torch.cuda.empty_cache()
return state_dict, num_samples

```

Appendix F — Development History

Previous Report: Literature review on federated learning and generative models for medical imaging; system architecture design; GAN-based prototype using the Flower framework; proof-of-concept results achieving 78.79% global classifier accuracy over 5 communication rounds.

Final Report: The project upgraded from GANs to a class-conditioned DDPM diffusion model for skin-lesion synthesis on the HAM10000 dataset, added DDIM fast sampling for efficient inference, replaced the Flower framework with a custom PyTorch FedAvg implementation, quantified synthetic quality via FID (raw malignant FID 218), and added a federated EfficientNet-B0 classifier. The final augmentation experiment raised malignant recall from 0.76 to 0.79 at 91.81% overall accuracy and 0.9559 AUC-ROC, confirming the end-to-end benefit of the pipeline.

Appendix G — Additional Synthetic Output Samples




Figure F.1: Additional synthetic malignant samples (federated DDIM, 64x64).

Random Sample — 16 of 1500 Synthetic Malignant Images (DDIM, 64x64)




Figure F.2: Additional synthetic malignant samples (federated DDIM, 64x64).




Figure F.3: Additional synthetic malignant samples (federated DDIM, 64x64).

```

.. =====
  BASELINE   (real data only)   : 91.66%
  AUGMENTED  (real + DDIM synth) : 91.81%
  Delta      : ▲ 0.15%
  =====

  Baseline AUC-ROC : 0.9572
  Augmented AUC-ROC : 0.9559
  Delta          : -0.0013
  
```

Figure F.4: Baseline vs. augmented classifier comparison summary.